

BRD — Versant Prep Tool (Business Requirements Document)

← Back to Index

Version 0.5, draft, not approved. Owner TBD. Companion doc: FRD — Versant Prep Tool. This note condenses the source BRD; see the original PDF for full appendix detail (v0.3 and earlier).

Changes since v0.3: product direction expanded to include a gamification/engagement layer and a dual B2C + B2B go-to-market, per Gamification & Dual-Market Expansion — Proposed Direction. New **Objective O7** (engagement without corrupting the readiness signal). **Risk R11** broadened from Coach mode specifically to any engagement feature. **Assumption A1** narrowed to the B2B motion specifically; new **Assumption A7** added for the B2C motion, unvalidated. New stakeholder: individual B2C learner. Scope (§7) extended to include the engagement layer and self-serve B2C signup in Phase 1. New **Risk R14** (gamification undermining B2B trust). New **Open Decisions D10–D13** for the mechanics still unresolved (badge visibility to cohort admins, B2C free-tier depth, whether gamification touches Exam-mode results screens, one vs. two onboarding flows). Commercials (§11) updated to name the dual pricing-model question explicitly. New Open Decisions D10-D13 were subsequently resolved in the same v0.4 cycle (see §12) and are now implemented as real requirements in FRD — Versant Prep Tool v0.5 §3.9-3.10 (FR-800 and FR-820 series).

Changes in v0.5: scope (§7) extended with the onboarding goal questionnaire, the Personalized Roadmap, and the full Admin/Super Admin/Org Admin role hierarchy (reserved/Phase 2), per Versant Tool — Phases by Module and Login, Roles, Permissions & Onboarding. Phasing (§10) clarified: build does not wait on A1, since A7 is independent of it; Coach Mode moved to a post-launch fast-follow. New Open Decisions D14-D16.

How to read this document: it describes a product whose **core market premise has not yet been validated** (Assumption A1 — the only assumption still blocking Phase 0 exit). Factual statements about the Versant test are drawn from Pearson documents dated 2021-2022 and were re-verified directly against Pearson's live sources on 2026-08-25 (Assumption A2, now closed). Claims sourced from secondary research that couldn't be verified are logged in Appendix B and must not be used as a basis for commitment.

1. Business context

What Versant is: a fully automated English assessment by Pearson, ~30 minutes, ~70 items across 6 task types. Reports an Overall score plus four skill scores (Speaking, Listening, Reading, Writing), each weighted equally at 25%, on Pearson's Global Scale of English (10-90), with a corresponding CEFR level. Full task structure: Versant Test — Official Task Structure.

How it's used: employer-administered hiring screen, not a portable credential. Candidates cannot self-book — a Test Identification Number is issued by a test administrator, results go to the purchasing organisation via Pearson's ScoreKeeper portal. Pearson does not set a pass mark; the purchasing organisation sets its own cut score.

Who sits it: predominantly candidates in recruitment pipelines for BPO, call-centre, and outsourced customer service roles, concentrated in India and the Philippines (scale unverified — see Appendix B).

What this means for the candidate: a high-stakes, short-notice, employer-controlled gate — one attempt, a result they may never see in full, failure blocks employment.

What this means commercially: the buyer of the test is the employer — the party motivated to improve outcomes may be the employer, a staffing agency, a training academy, or the candidate, and these have very different willingness/ability to pay. Resolving this is **Assumption A1**.

2. Problem statement

For the candidate: existing preparation is either a fixed-form paid practice test that returns a score with no diagnosis, or unstructured free content. Neither answers *am I ready?* or *what specifically should I fix?* The test measures automaticity under time pressure, which reading explanations cannot improve — it requires repeated timed production with corrective feedback.

For an organisation preparing candidates: no instrument exists to identify which candidates are ready, which are close and worth training, and which are not viable — other than paying to put everyone through the real test and reading the failures afterward.

3. Objectives and success criteria

ID	Objective	Success criterion
----	-----------	-------------------

ID	Objective	Success criterion
O1	Give a learner an accurate readiness signal	Readiness band agrees with actual pass/fail outcome (vs. the client's cut score) for $\geq 75\%$ of a validation cohort of ≥ 100 learners
O2	Give feedback a learner can act on	$\geq 60\%$ of learners who complete 3+ practice sessions show measurable improvement on their weakest task
O3	Faithfully reproduce test conditions	Item formats, word-length ranges, vocabulary bands and timing windows match published Pearson specifications, verified by review against Appendix A
O4	Score speech fairly across accents	Content-scoring disagreement with human transcription $\leq 10\%$ word error on a held-out set spanning ≥ 4 L1 backgrounds
O5	Establish a defensible commercial line	At least one paying design-partner cohort by end of Phase 2
O6	Coach mode feels conversational, not laggy	Median agent turn latency under 1s; speech onset detection accurate within 100ms of hand-marked onset on a test set
O7	Engagement/gamification layer increases practice frequency without corrupting the readiness signal	Streak/badge engagement metrics and readiness-band accuracy (O1) tracked as two separate series; ship only features that move practice frequency and show no negative correlation with O1's validation accuracy

O1 and O4 are the objectives product credibility rests on — if not met, the product should not ship regardless of the state of the rest. **O7 is explicitly not allowed to trade against O1** — see Risk R11.

4. Stakeholders

Stakeholder	Interest
Learner / candidate	Passing the screen; knowing what to practise
Individual B2C learner (student / job-seeker, self-pay)	A product that's worth paying for on its own merits — engaging enough to return to daily, not just a hiring-mandated chore. New stakeholder as of v0.4; not present when the product was B2B-only
Cohort administrator (academy trainer, L&D or TA coordinator)	Cohort readiness visibility; pass-rate improvement
Purchasing organisation	Reduced hiring cost per qualified candidate
Edstellar sales / account management	A sellable proposition into an identified buyer
Edstellar delivery / L&D	Integration with existing training offerings
Engineering	Buildability; scoring accuracy; unit cost
Legal / compliance	Trademark exposure; voice data as sensitive personal data

(This confirms Edstellar — an existing training/L&D company — is building this as a diversification, not a startup from scratch.)

5. What the "AI" actually is

Written for non-technical readers, since "AI tutor" hides four different technologies doing four different jobs with very different importance:

The two that are the product:

- **Speech recognition** — determines which words were produced and exactly when, constrained against the known target text (not guessing freely) — this is what keeps scoring fair for strong accents.
- **Pronunciation and fluency measurement** — from word/sound timings, derives which sounds were wrong, where speech broke up, speaking rate, latency. A measurement model. Outputs numbers, not sentences.

The two that are supporting cast:

- **A language model** — used in exactly three places: drafting practice items (human-reviewed before release), turning measurement output into readable feedback, and providing Coach mode's voice.
- **Speech synthesis** — generates item audio and Coach mode's voice.

What the language model does not do: it does not decide anything, does not judge whether a learner got an item right, does not assess pronunciation, does not rephrase what the measurement model already determined. A language model asked to judge pronunciation from a transcript would produce confident, fluent, *incorrect* judgements — and teach the learner something false. **The scorer decides; the language model narrates.** This is the reasoning behind FRD constraint C11 and requirement FR-704.

Where there is deliberately no AI at all:

Function	How it actually works
Choosing what to practise next	Database query: weakest task, most frequent error patterns
Enforcing response windows	A server clock
Readiness band	A lookup table against configured thresholds
Item difficulty ordering	Stored difficulty estimates from observed performance

What could become genuinely proprietary: in Phase 1 the measurement models are bought/borrowed. With enough learner recordings and enough paired real test outcomes, they could eventually be retrained — better tuned to Indian and Filipino accents, calibrated against actual Versant results rather than a published CEFR mapping. That depends on the calibration data chain: A1 (a reachable buyer) → client cohorts → A4 (paired real outcomes).

6. Assumptions and dependencies

ID	Assumption	Status	Evidence required
A1	(B2B motion) There is a reachable employer/agency/academy buyer for Versant preparation within or adjacent to Edstellar's existing client relationships	UNVALIDATED — highest risk in this document	CRM query for accounts administering Versant/English screening; 8-10 structured conversations with existing clients; ≥2 unprompted mentions of Versant or an equivalent screen. External signal found 2026-08-25 (TDS Global Solutions actively sells Versant screening to BPOs) — external evidence only, does not substitute for internal CRM/client-conversation evidence. Narrowed to "B2B motion" specifically as of v0.4, now that A7 covers the B2C motion separately
A2	Test format, item counts and timings remain as documented in 2021-2022	VERIFIED 2026-08-25	Re-fetched both source PDFs directly from Pearson's live URLs and extracted their real text (not model-summarized — a first-pass model summary of the validation PDF was caught fabricating a different, wrong table and discarded). Confirmed word-for-word: task structure (16/8/12/18/14/2 items, Parts A-F), all five response windows (15s/6s, 15s/8s, 25s, 25s, 30s+90s), word-count ranges (3-15, 3-14, 45-65), and the exact GSE→CEFR Table 7 (20-23 <A1 ... 79-80 C2) all match Appendix A precisely. Both source documents are unchanged since their original dates (April 2021 v1.0 guide; 2022 validation summary) despite being fetched fresh today.
A3	No incumbent AI-based Versant preparation product holds the market	FALSE — resolved 2026-08-25	Direct search and hands-on evaluation found ≥14 Versant-specific sites/apps (VersantPrep, Mocko, VERSANT Bridge, CA Monk, MockVersant, VersantPro, MyEnglishPractice.jp, etc.) plus Pearson's own channel. Every one is an individual-learner, self-pay product; none serve an employer-facing B2B buyer with a compliance pipeline. Per Risk R6, this resolves as evaluate and reposition, not stop — see Market Gap & Positioning
A4	Real Versant scores can eventually be obtained for calibration	Unvalidated	Depends on A1 — without client cohorts, calibration relies on slower, noisier self-reported scores from individual users
A5	Published Versant→CEFR mapping is a valid anchor for a readiness band	Partially supported	Pearson's own validation summary provides the mapping (see Versant Test — Official Task Structure); its applicability to this product's independent measures is untested
A6	Learners will accept a readiness band rather than a numeric score	Unvalidated	Test in Phase 1 with real users
A7	(B2C motion, new v0.4) Individual candidates/students will pay for this directly (self-serve), independent of any employer relationship	UNVALIDATED — new, no evidence gathered yet	Requires its own validation track, separate from A1: pricing sensitivity testing against the individual competitors in AI Speaking Apps — Landscape and Versant Sites — Quick Reference Directory (the B2C competitive bar is far more crowded than the B2B one — see Market Gap & Positioning); a real free-tier design that doesn't repeat TalkPal's "free in name only" mistake

Dependency: Phase 3 calibration is blocked on A1 and A4. **A7 is independent of A1** — the B2C motion can proceed or fail without affecting the B2B motion's validation, and vice versa.

7. Scope

In scope — Phase 1

- Learner-facing web application, mobile-responsive
- Three of the six task types: **Repeats, Sentence Builds, Dictation**
- **Exam mode** — silent, deterministic, timed item delivery, no assistance, server-enforced, matching published response windows. The only

environment scores are produced in.

- **Coach mode** — a conversational voice agent that drills bank items, corrects pronunciation aloud, and debriefs a completed session. Nothing here is scored.
- A hard boundary between the two modes (Constraint C9)
- Automated scoring: content accuracy, fluency, pronunciation, mechanical written accuracy
- Per-task diagnostic feedback
- Readiness band, expressed as a band with per-task detail — **not** a numeric Versant score
- Original item bank generated against published specifications, human-reviewed
- Consent, retention and deletion handling for voice data
- Multi-tenant data model (no cohort UI yet) — now doubling as the B2C/B2B split: an individual B2C learner is simply a learner on the default tenant, no separate build needed
- **(New, v0.4)** A non-scored engagement/gamification layer (streaks, accuracy-gated leveling, achievement badges) — lives entirely outside Exam mode, per O7 and the extended R11. Learner-private only (D10); confined to Coach mode and the dashboard, never the Exam-mode results screen (D12)
- **(New, v0.4)** Self-serve B2C signup path — one shared flow branching by entry point (D13), with a real, volume-capped (not feature-gated) free tier (D11)
- **(New, v0.5)** Onboarding **goal questionnaire** — exam choice, target CEFR level, exam date, prep days, daily minutes, past attempts — collected before the placement test
- **(New, v0.5)** The **"you may already be at your target level"** check against the stated goal
- **(New, v0.5)** **Personalized Roadmap** — rule-based study plan with milestones, from placement result + goal + exam date
- **(New, v0.5)** **Super Admin / Admin / Org Admin role hierarchy**, per Login, Roles, Permissions & Onboarding — reserved, not built. No separate reviewer role; content/score review are Admin permissions

Open Decisions D10-D13 are now resolved — see §12 and FRD — Versant Prep Tool v0.5 §3.9-3.10 for the corresponding requirements.

Within Phase 1, **Exam mode ships before Coach mode** — Exam mode is the product, Coach mode makes it habitable. Building the agent first risks an engaging practice companion that doesn't move scores (Risk R11).

Explicitly out of scope — Phase 1

Reading skill coverage (the two tasks that carry Reading — Sentence Completion, Passage Reconstruction — are deferred, so Phase 1 covers three of four skills and readiness output must state this limitation); cohort administrator dashboards/reporting; numeric predicted Versant score; other Versant product variants; native mobile applications; integration with client HR/ATS systems.

Deferred — later phases

Item	Phase	Reason for deferral
Sentence Completion	2	Technically simple, but requires curated answer sets with collocation preferences for 18 items per test — a content-production problem
Conversations	2	Requires semantic scoring of open-ended short spoken answers
Passage Reconstruction	3	Hardest scorer: semantic content coverage plus mechanical accuracy
Cohort admin and reporting	2	Required as soon as an organisational buyer is confirmed (A1)
Numeric score prediction	3	Blocked on calibration data (A4)

8. Constraints

ID	Constraint
C1	No reproduction of Pearson test items, audio, or copyrighted rubric text. All items original, generated against published specifications.
C2	No claim of affiliation, endorsement, or official status. "Official," "endorsed," and Pearson logos are not permitted. Legal review required before any public-facing copy.
C3	Voice recordings are personal data and may be biometric data. Explicit opt-in consent, stated retention limits, deletion on request are mandatory — informed by GDPR, India's DPDPA 2023, and the Philippines Data Privacy Act (see Appendix C).
C4	Item vocabulary must respect published frequency bands: spoken items to the 5,000 most frequent Switchboard Corpus words; written items to the 1,600 most frequent Longman Corpus Network words.

ID	Constraint
C5	Item topics restricted to the published workplace domain list (announcements, business trips, complaints, customer service, telephone/email, inventory, scheduling, marketing/sales).
C6	Items must be context-independent — no dependence on world knowledge or prior work experience.
C7	Third-party ASR or scoring services require a data processing agreement with provider-side audio retention disabled.
C8	Target infrastructure cost per full mock session to be set as a hard budget in the FRD; scoring architecture chosen to meet it.
C9	Exam mode and Coach mode are separated absolutely. Exam mode offers no assistance, no repeats, no encouragement, no agent presence. Nothing produced in Coach mode contributes to a score. A practice environment kinder than the real test teaches the learner nothing about the real test.
C10	A headset with a boom microphone is required — a hard requirement, not a recommendation. Matches the real test's stated requirement, removes the need for far-field processing/echo cancellation, materially improves scoring accuracy.
C11	The Coach mode agent shall not improvise language corrections. Any pronunciation or accuracy correction it speaks must originate from scorer output, not the language model's own judgement.
C12	Scored measurement is derived only from full-utterance constrained alignment. Streaming partial recognition may drive interface feedback but shall never contribute to a score.

9. Risks

ID	Risk	Severity	Mitigation	Owner
R1	A1 is false: no reachable buyer	Critical	Validate before build spend. Fallback: reframe as a general business-English assessment sold into existing L&D relationships, with Versant as one item-format module	Product / Sales
R2	Readiness signal is inaccurate (O1 fails)	Critical	Conservative banding; state uncertainty; withhold numeric claims until calibrated; collect paired outcomes from first cohorts	Data / Product
R3	Speech scoring penalises accents (O4 fails)	High	Constrained recognition against known target text; confidence thresholds; human-review path for disputed items; test across ≥4 L1 backgrounds	Engineering
R4	Trademark or content dispute	High	C1 and C2; legal review of naming and marketing copy before launch	Legal
R5	Voice data compliance failure	High	C3, C7; privacy design reviewed before first external user	Legal / Engineering
R6	An incumbent product already owns the niche (A3 false)	Medium — CLOSED 2026-08-25	Verified: real incumbents exist, but every one is an individual-learner, self-pay product; none serve an employer-facing B2B buyer with a compliance pipeline. Resolution: reposition around B2B/compliance-rigorous delivery, not stop	Product
R7	Test specification has changed since 2022 (A2 false)	Medium — CLOSED 2026-08-25	Re-verified directly against Pearson's live source PDFs (see A2). No drift found — the 2021/2022 documents are still the current ones served at the same URLs. Re-check again immediately before item generation begins regardless, since this can change without notice	Product
R8	Free content is good enough for learners	Medium	Differentiate on diagnosis and timed repetition, not content volume	Product
R9	Unit cost per session exceeds price point	Medium	Cost model built before scoring vendor selection; prototype on hosted API, plan migration path to self-hosted	Engineering
R10	Coach mode cost dominates unit economics	High	Hard per-learner daily cap on agent minutes; default to structured drill over open conversation; measure cost per learner from week one	Engineering / Product
R11	(Broadened v0.4) Engagement features — Coach mode <i>and</i> the gamification layer (streaks, badges, leveling) — raise usage but not scores	High	Coach mode delivers bank items under real timing, not improvised conversation; track score movement separately from all engagement metrics (O7); cut features that lift only the latter. Leveling must be accuracy-gated, not time-on-task-gated — the specific mistake found in TalkPal — UX Engagement Analysis	Product
R12	Onset detection failure — a missed onset ends the item and produces a wrong score	High	Tune asymmetrically toward early triggering; verify against the recording in the batch pass; allow the learner to flag the item	Engineering
R13	Mode bleed — assistance/retries/agent presence leaking into Exam mode via well-intentioned UX decisions	Medium	C9 stated as a constraint; mode separation asserted as a testable requirement, not a design convention	Product / Engineering

ID	Risk	Severity	Mitigation	Owner
R14	(New, v0.4) Gamification undermines B2B trust — a skeptical employer buyer sees streaks/leaderboards/celebratory copy as gimmicky or, worse, as surveillance of their candidates	Medium	Keep badge/streak visibility learner-private by default (Open Decision D10); B2B-facing surfaces (cohort dashboards, certificates) stay in the sober, diagnostic register the BRD already uses elsewhere — VersantPro's fake "students practicing now" counter is the cautionary example (see What To Take — Design Lessons)	Product

10. Phasing

Versant Tool — Phases by Module is the granular engineering build order (Phases 0-11) underlying Phase 1 below.

- **Phase 0 — Validate and specify** (target 2-3 weeks): confirm calibration dataset licensing; legal read on naming/claims. A2 and A3 both closed 2026-08-25. **A1 no longer gates engineering start** (v0.5) — Assumption A7 (B2C) is independent of A1, so Foundation/build work starts immediately; A1 evidence-gathering runs in parallel and gates Phase 2's B2B/cohort work instead.
- **Phase 1 — Learner MVP**: three tasks, scoring engine (placeholder to start, real vendor mid-build — FRD FR-245), diagnostics, goal questionnaire, Personalized Roadmap, readiness band, consent handling, non-scored engagement layer. Coach Mode is a **post-launch fast-follow** (v0.5), not in the launch scope. **Exit**: O3 and O4 met, ≥20 real learners completed a full session.
- **Phase 2 — Cohort layer and scope extension**: Sentence Completion and Conversations; Org Admin/Admin/Super Admin screens (design complete — Login, Roles, Permissions & Onboarding); cohort administrator reporting; first paying design-partner cohort. **Exit**: O5 met, ≥100 learners with recorded outcomes, A1 evidenced.
- **Phase 3 — Calibration and completion**: Passage Reconstruction; calibration against paired real Versant outcomes; numeric score prediction only if accuracy supports it. **Exit**: O1 met.

16-Week Plan (2026-08-27): Phase 1 runs as 3 parallel tracks over 16 weeks (3 people); AI vendor decision is a hard week-3 deadline, no second round budgeted (hardens D3/FRD TD-01).

11. Commercial

Deliberately left open. Pricing model (per-seat, per-cohort licence, or subscription) depends entirely on which buyer A1 identifies — setting it before that is guesswork. Cost inputs (per-session infrastructure cost) established in Phase 0 so a price floor is known.

As of v0.4: this is now explicitly a **dual pricing question**, not a single one — an individual B2C subscription tier (A7) and a B2B per-seat/per-cohort licence (A1) most likely coexist rather than one replacing the other. VERSANT Bridge — Competitor Deep-Dive is the closest working reference for this exact dual structure — individual subscriptions and corporate certificates sold side by side. Still fully open: whether the B2C tier includes a real free plan (see A7's evidence requirement) and how the B2B tier's certificate/report value is priced relative to it.

12. Open decisions

ID	Decision	Needed by
D1	Confirm or reframe A1	Phase 0 exit — gates everything
D2	Product name (constrained by C2)	Before any public copy
D3	Speech scoring: hosted API vs self-hosted stack	Direction set : start hosted, behind a swappable interface (FRD TD-01). Specific vendor open.
D4	Voice retention period and default	Resolved : 90 days default from capture, configurable per tenant (FR-601)
D5	Whether Phase 1 is offered free to gather calibration data	Phase 1 planning
D6	Number of readiness bands and their labels	Phase 1 design
D7	Realtime speech-to-speech provider for Coach mode	Before Coach mode build
D8	Agent minute cap per learner per day, and whether Coach mode is gated behind a paid tier	Before Coach mode build
D9	Whether headsets are enforced by blocking the session or by warning only (C10)	Phase 1 design
D10	Resolved . Badges/streaks are strictly private to the learner — no cohort administrator view ever displays individual gamification data. A cohort administrator instead sees aggregate, non-comparative practice-frequency stats (Phase 2).	FRD FR-825/FR-826

ID	Decision	Needed by
D11	Resolved. A real B2C free tier exists, capped by volume (daily/weekly scored-attempt and Coach-minute limits) not by feature/task-type gating — avoids TalkPal's "one mode only" mistake.	FRD FR-802-FR-805
D12	Resolved. Gamification never touches the Exam-mode session-completion/readiness screen — that stays strictly diagnostic. Confined to Coach mode and a separate dashboard/history view.	FRD FR-828
D13	Resolved. One shared signup flow, branching by entry point (public link vs. cohort invitation token) — not two separate onboarding builds.	FRD FR-801
D14	Resolved, v0.5. Account model is one <code>users</code> table with <code>role ∈ {learner, org_admin, admin, super_admin}</code> . No separate reviewer account type — <code>content_review / score_review</code> are permissions an Admin can hold. Admin capabilities are individually permission-gated by a Super Admin, not fixed per department label. Full spec: Login, Roles, Permissions & Onboarding.	FRD §3.12 (FR-950 series)
D15	Resolved, v0.5. The AI speech-scoring vendor decision (D3/TD-01) is scheduled for week 3 of the Phase 1 build, with no second evaluation round budgeted — a hard commitment, not just a direction.	16-Week Plan, week 3
D16	Resolved, v0.5. No content-reviewer UI is built in Phase 1. Approved items are inserted directly into the item bank by the team, pre-approved; FR-403/404's reviewer fields remain in the schema, populated manually, for the day a real review workflow is built.	FRD FR-403 (Phase 1 note)

Appendix A — Verified test facts

See Versant Test — Official Task Structure for the full extracted task structure, timing, scoring approach, and GSE↔CEFR table, sourced from Pearson's *Versant 4 Skills Essential Test: Official Guide for Test-Takers* (2021) and *Test Description and Validation Summary* (2022). Must be re-verified per Assumption A2.

Appendix B — Unverified claims register

Claims encountered in secondary research that could **not** be verified and must not be relied upon:

Claim	Status
An AI-driven Versant prep platform exists covering all six parts at ~\$5/month, with "thousands" of users	Partially confirmed 2026-08-25 — the general shape is now true (multiple real platforms exist, one covers 9 parts of one variant, several are free/low-cost), but the specific figures remain unverified for any single product
Pearson delivers ~3 million Versant tests annually	Unverified
Philippine 4-Skills mean score ≈63; local cut scores mid-50s to 60s	Unverified
GSE 59 ≈ Versant 58, and CEFR B2 = Versant 58-68	Rejected — contradicts primary Pearson documents (GSE 59/90 = Versant 57/80; B2 = 57-67). Use Appendix A.
Pearson official practice test priced at \$15	Unverified
Versant product variants and their durations (Placement 50 min, Professional 60 min, Speaking & Listening 20 min, Writing 35 min)	Plausible, unverified

Appendix C — Regulatory frameworks

Expands constraints C3 and C7. Background for legal review, not a substitute for it.

Framework	Region	Why relevant
GDPR	EU/EEA	Applies to any learner located in the EU/EEA regardless of company location; voice recordings may be "special category" biometric data
DPDPA 2023	India	Directly relevant — India is a named data-residency market and a concentration of the target candidate population
Data Privacy Act of 2012 (RA 10173)	Philippines	Directly relevant — the other named residency market; National Privacy Commission guidance treats biometric data as sensitive personal information

Common principles across all three: consent must be specific and affirmative (FR-011); purpose limitation — data can't be reused for a new purpose like model training without separate consent (FR-603); storage/retention limitation (FR-601, full matrix in FRD §6); right to erasure, propagating across every system (FR-602); processor accountability — any third party needs a contractual DPA (FR-605, C7); access control and auditability (FR-606).

Where they diverge: GDPR is restrictive-by-default on cross-border transfer (approved mechanism required); DPDPA is permissive-by-default (transfers allowed except to blacklisted countries); Philippines requires case-by-case adequate safeguards. Maximum penalties range from up to 4% of global annual revenue (GDPR) to ~\$30M (DPDPA) to RA 10173's National Privacy Commission-enforced penalties. None of the "extra obligations at scale" (DPO requirements, Significant Data Fiduciary designation) are Phase 1 blockers, but should be revisited with counsel once real user volume exists.